

Introduction to Deep Learning for Computer Vision



Adhyayan '23 - ACA Summer School
Department of Computer Science and Engineering
Indian Institute of Technology Kanpur

Lecture 2

Training Neural Networks!

Loss Optimization

- We want to find network weights that *achieve the lowest loss*.

$$\mathbf{W}^* = \operatorname{argmin}_{\mathbf{W}} \frac{1}{n} \sum_{i=1}^n \mathcal{L}(f(x^{(i)}; \mathbf{W}), y^{(i)})$$

$$\mathbf{W}^* = \operatorname{argmin}_{\mathbf{W}} J(\mathbf{W})$$


Remember:

$$\mathbf{W} = \{W^{(0)}, W^{(1)}, \dots\}$$

Loss Optimization


$$\mathbf{W}^* = \operatorname{argmin}_{\mathbf{W}} J(\mathbf{W})$$

Remember:
*Our loss is a function of
the network weights.*




Loss Optimization

Randomly pick an initial (w_0, w_1)




Loss Optimization

Compute gradient, $\frac{\partial J(W)}{\partial W}$




Loss Optimization

Take small step in opposite direction of the gradient



Gradient Descent


Repeat Until Convergence



Gradient Descent


- Algorithm:
 1. Initialize weights randomly $\sim \mathcal{N}(0, \sigma^2)$
 2. Loop until convergence:
 3. Compute gradient, $\frac{\partial J(W)}{\partial W}$
 4. Update weights, $W \leftarrow W - \eta \frac{\partial J(W)}{\partial W}$
 5. Return weights

Computing Gradients: Backpropagation



How does a small change in one weight (ex. w_2) affect the final loss $J(\mathbf{W})$


Computing Gradients: Backpropagation



$$\frac{\partial J(W)}{\partial w_2}$$


Let's use chain rule!

Computing Gradients: Backpropagation



$$\frac{\partial J(W)}{\partial w_2} = \underbrace{\frac{\partial J(w)}{\partial \hat{y}}}_{\text{purple}} \times \underbrace{\frac{\partial \hat{y}}{\partial w_2}}_{\text{red}}$$

Computing Gradients: Backpropagation




$$\frac{\partial J(W)}{\partial w_1} = \frac{\partial J(w)}{\partial \hat{y}} \times \frac{\partial \hat{y}}{\partial w_1}$$

Apply Chain Rule


Apply Chain Rule

Computing Gradients: Backpropagation



$$\frac{\partial J(W)}{\partial w_1} = \underbrace{\frac{\partial J(w)}{\partial \hat{y}}}_{\text{purple}} \times \underbrace{\frac{\partial \hat{y}}{\partial z_1}}_{\text{red}} \times \underbrace{\frac{\partial z_1}{\partial w_1}}_{\text{blue}}$$

Computing Gradients: Backpropagation



$$\frac{\partial J(W)}{\partial w_1} = \underbrace{\frac{\partial J(w)}{\partial \hat{y}}}_{\text{red}} \times \underbrace{\frac{\partial \hat{y}}{\partial z_1}}_{\text{red}} \times \underbrace{\frac{\partial z_1}{\partial w_1}}_{\text{blue}}$$

Repeat this for **every weight in the network** using gradients from later layers

Neural Networks in Practice: Optimization!

Training Neural Networks is Difficult!




Image: Li, Hao, et al. "Visualizing the loss landscape of neural nets." Advances in neural information processing systems 31 (2018).

Loss Functions Can Be Difficult To Optimize

Remember: Optimization through
Gradient Descent

$$W \leftarrow W - \eta \frac{\partial J(W)}{\partial W}$$

Loss Functions Can Be Difficult To Optimize


Remember: Optimization through
Gradient Descent

$$W \leftarrow W - \eta \frac{\partial J(W)}{\partial W}$$

How can we set the
learning rate?


Setting the Learning Rate

Small learning rates converge slowly and get stuck in a false local minima.




Setting the Learning Rate

Large learning rates overshoot, become unstable and diverge.



Setting the Learning Rate

Stable learning rates converge smoothly and avoid local minima.



How to deal with this?

Idea 1:

Try lots of different learning rates and see what works “just right”.

How to deal with this?

Idea 1:

Try lots of different learning rates and see what works “just right”.

Idea 2:

Do something smarter!

Design an adaptive learning rate that “adapts” to the landscape!

Adaptive Learning Rates

- Learning Rates are no longer fixed
- Can be made larger or smaller depending on:
 - How large gradient is
 - How fast learning is happening
 - Size of particular weights
 - etc...

Gradient Descent Algorithms

Algorithm	Reference
SGD	Kiefer, Jack, and Jacob Wolfowitz. "Stochastic estimation of the maximum of a regression function." <i>The Annals of Mathematical Statistics</i> (1952): 462-466.
Adam	Kingma, Diederik P., and Jimmy Ba. "Adam: A method for stochastic optimization." <i>arXiv preprint arXiv:1412.6980</i> (2014).
Adadelata	Zeiler, Matthew D. "Adadelata: an adaptive learning rate method." <i>arXiv preprint arXiv:1212.5701</i> (2012).
Adagrad	Duchi, John, Elad Hazan, and Yoram Singer. "Adaptive subgradient methods for online learning and stochastic optimization." <i>Journal of machine learning research</i> 12.7 (2011).
RMSProp	Tieleman, T., & Hinton, G. (2012). Lecture 6.5-rmsprop: divide the gradient by a running average of its recent magnitude. COURSERA: Neural networks for machine learning, 4(2), 26–31.


Neural Networks in Practice: Mini-batches!

Gradient Descent

- Algorithm:

1. Initialize weights randomly $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Compute gradient, $\frac{\partial J(W)}{\partial W}$
4. Update weights, $W \leftarrow W - \eta \frac{\partial J(W)}{\partial W}$
5. Return weights

$J(w_0, w_1)$




Can be very
computationally
intensive to compute.

Stochastic Gradient Descent

- Algorithm:

1. Initialize weights randomly $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Pick a single data point \mathbf{i}
4. Compute gradient, $\frac{\partial J_i(W)}{\partial W}$
5. Update weights, $W \leftarrow W - \eta \frac{\partial J(W)}{\partial W}$
6. Return weights

$$J(w_0, w_1)$$




Easy to compute but
very noisy (stochastic)!

Stochastic Gradient Descent

- Algorithm:

1. Initialize weights randomly $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Pick a batch of **B** data points.
4. Compute gradient, $\frac{\partial J(W)}{\partial W} = \frac{1}{B} \sum_{k=1}^B \frac{\partial J_k(W)}{\partial W}$
5. Update weights, $W \leftarrow W - \eta \frac{\partial J(W)}{\partial W}$
6. Return weights




Fast to compute and a much better estimate of the true gradient!

Mini-batches while training

- More accurate estimation of gradient:
 - Smoother convergence
 - Allows for larger learning rate
- Mini-batches lead to fast training!
 - Can parallelize computation
 - Achieve significant speed increases of GPUs!

Neural Networks in Practice: Overfitting!

The Problem of Overfitting



Model does not have capacity
to fully learn the data

Too complex, extra parameters
do not generalize well

Regularization


- What is it?
 - Technique that constrains our optimization problem to discourage complex models

Regularization


- What is it?
 - Technique that constrains our optimization problem to discourage complex models
- Why do we need it?
 - Improve generalization of our model on unseen data.

Regularization I : Dropout!

- During training, randomly set some activations to 0.
 - Typically 'drop' 50% of activations in layer
 - Forces the network to not rely on any one node.




(a) Standard Neural Net



(b) After applying dropout.

Regularization II: Early Stopping

- Stop training before we have a possibility to overfit.



Next Lecture: Convolutional Neural Networks!