


# Introduction to Deep Learning for Computer Vision

---

Adhyayan '23 - ACA Summer School  
Department of Computer Science and Engineering  
Indian Institute of Technology Kanpur

Lecture 3

# Tasks in Computer Vision



- **Regression:** Output variable takes continuous values.
- **Classification:** Output variable takes class label. Can produce probability of belonging to a particular class.

# High Level Feature Detection

Let's identify key-features in each image category



Nose,  
Eyes,  
Mouth




Wheels,  
License Plate,  
Headlights



Doors,  
Windows,  
Steps

# Learning Feature Representations

Can we learn a **hierarchy of features** directly from data instead of hand engineering?




Edges, dark spots

Eyes, ears, nose

Facial Structure


# **Learning Visual Features**

# Fully Connected Neural Network



# Fully Connected Neural Network

- Input:
  - 2D Image
  - Vector of pixel values




- Fully Connected:
  - Connect neuron in hidden layer to all neurons in output layer
  - No spatial information!
  - And many, many parameters.

# Fully Connected Neural Network

- Input:

- 2D Image
- Vector of pixel values



- Fully Connected:


- Connect neuron in hidden layer to all neurons in output layer
- No spatial information!
- And many, many parameters.

How can we use **spatial structure** of the input to inform the architecture of the neural network?





# Using Spatial Structure

- Input:
  - 2D Image
  - Array of pixel values



Idea: Connect patches of input to neurons in hidden layer. Neurons connected to region of input, only 'sees' these values.

# Using Spatial Structure



Connect patch in input layer to single neuron in subsequent layer.


Using a sliding window to define connections.

*How can we **weight** the patch to detect particular features?*

# Applying Filters to Extract Features

1. Apply a set of weights - a filter - to extract **local features**
2. Use multiple filters to extract different features.
3. Spatially share parameters of each filter.  
(features that matter in one part of the input should matter elsewhere)

# Feature extraction with Convolution



- Filter of size 4 x 4: 16 different weights
- Apply the same filter to 4 x 4 patches in input
- Shift by 2 pixels for next patch

This 'patchy' operation is called **convolution**

1. Apply a set of weights - a filter - to extract local features
2. Use multiple filters to extract different features.
3. Spatially share parameters of each filter.  
(features that matter in one part of the input should matter elsewhere)

# **Feature Extraction and Convolution: A Case Study**

# X or X?


-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1



-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	1	-1	-1	-1
-1	1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	1	1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	1	-1	-1	-1
-1	-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1

Image is represented as matrix of pixel values...and computers are literal!  
We want X to be classified as X even if it is shifted, rotated, shrunk or deformed.

# Features of X




# Filters to Detect X Features

filters

1	-1	-1
-1	1	-1
-1	-1	1


1	-1	1
-1	1	-1
1	-1	1

-1	-1	1
-1	1	-1
1	-1	-1






# The Convolution Operation



# The Convolution Operation

Suppose we want to compute the convolution of a 5x5 image and a 3x3 filter:

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0



1	0	1
0	1	0
1	0	1

We slide the 3x3 filter over the input image, element-wise multiply, and add the outputs...

# The Convolution Operation

We slide the 3x3 filter over the input image, element-wise multiply, and add the outputs:

1 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	0	0
0 <sub>x0</sub>	1 <sub>x1</sub>	1 <sub>x0</sub>	1	0
0 <sub>x1</sub>	0 <sub>x0</sub>	1 <sub>x1</sub>	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved  
Feature

# The Convolution Operation

We slide the 3x3 filter over the input image, element-wise multiply, and add the outputs:

1	1 <sub>x1</sub>	1 <sub>x0</sub>	0 <sub>x1</sub>	0
0	1 <sub>x0</sub>	1 <sub>x1</sub>	1 <sub>x0</sub>	0
0	0 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	1
0	0	1	1	0
0	1	1	0	0

Image

4	3	

Convolved  
Feature

# The Convolution Operation

We slide the 3x3 filter over the input image, element-wise multiply, and add the outputs:

1	1	1 <sub>x1</sub>	0 <sub>x0</sub>	0 <sub>x1</sub>
0	1	1 <sub>x0</sub>	1 <sub>x1</sub>	0 <sub>x0</sub>
0	0	1 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>
0	0	1	1	0
0	1	1	0	0

Image

4	3	4

Convolved  
Feature

# The Convolution Operation

We slide the 3x3 filter over the input image, element-wise multiply, and add the outputs:

1	1	1	0	0
0 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	1	0
0 <sub>x0</sub>	0 <sub>x1</sub>	1 <sub>x0</sub>	1	1
0 <sub>x1</sub>	0 <sub>x0</sub>	1 <sub>x1</sub>	1	0
0	1	1	0	0

Image

4	3	4
2		

Convolved  
Feature

# The Convolution Operation

We slide the 3x3 filter over the input image, element-wise multiply, and add the outputs:

1	1	1	0	0
0	1 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	0
0	0 <sub>x0</sub>	1 <sub>x1</sub>	1 <sub>x0</sub>	1
0	0 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	0
0	1	1	0	0

Image

4	3	4
2	4	

Convolved  
Feature

# The Convolution Operation

We slide the 3x3 filter over the input image, element-wise multiply, and add the outputs:

1	1	1	0	0
0	1	1 <sub>x1</sub>	1 <sub>x0</sub>	0 <sub>x1</sub>
0	0	1 <sub>x0</sub>	1 <sub>x1</sub>	1 <sub>x0</sub>
0	0	1 <sub>x1</sub>	1 <sub>x0</sub>	0 <sub>x1</sub>
0	1	1	0	0

Image

4	3	4
2	4	3

Convolved  
Feature



# The Convolution Operation

We slide the 3x3 filter over the input image, element-wise multiply, and add the outputs:

1	1	1	0	0
0	1	1	1	0
0 <sub>x1</sub>	0 <sub>x0</sub>	1 <sub>x1</sub>	1	1
0 <sub>x0</sub>	0 <sub>x1</sub>	1 <sub>x0</sub>	1	0
0 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	0	0

Image

4	3	4
2	4	3
2		

Convolved  
Feature

# The Convolution Operation

We slide the 3x3 filter over the input image, element-wise multiply, and add the outputs:

1	1	1	0	0
0	1	1	1	0
0	0 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	1
0	0 <sub>x0</sub>	1 <sub>x1</sub>	1 <sub>x0</sub>	0
0	1 <sub>x1</sub>	1 <sub>x0</sub>	0 <sub>x1</sub>	0

Image

4	3	4
2	4	3
2	3	

Convolved  
Feature

# The Convolution Operation

We slide the 3x3 filter over the input image, element-wise multiply, and add the outputs:


1	1	1	0	0
0	1	1	1	0
0	0	1 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x2</sub>
0	0	1 <sub>x0</sub>	1 <sub>x1</sub>	0 <sub>x0</sub>
0	1	1 <sub>x1</sub>	0 <sub>x0</sub>	0 <sub>x1</sub>

Image

4	3	4
2	4	3
2	3	4

Convolved  
Feature

# Producing Feature Maps



Original



Sharpen




Edge Detect



"Strong" Edge  
Detect

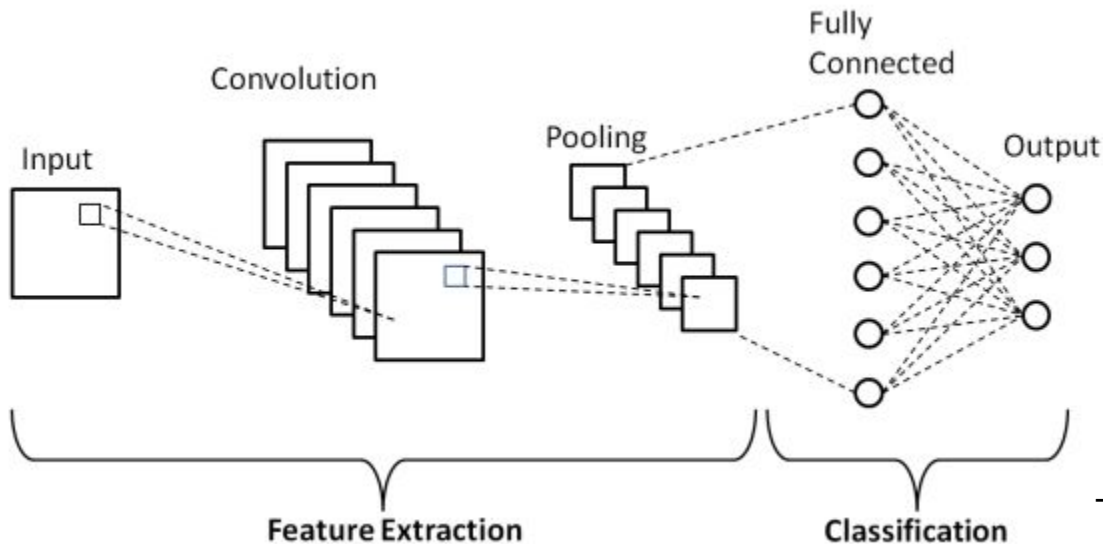
# Feature Extraction with Convolution



1. Apply a set of weights - a filter - to extract **local features**
2. Use multiple filters to extract different features.
3. Spatially share parameters of each filter.  
(features that matter in one part of the input should matter elsewhere)

# **Convolutional Neural Networks (CNNs)**


# CNNs for Classification



1. **Convolution:** Apply filters to generate feature maps.
2. **Non-Linearity:** Often ReLU.
3. **Pooling:** Downsampling operation on each feature map.

Train model with image data. Learn weights of filters in convolutional layers.

# Convolutional Layers: Local Connectivity



4x4 filter: matrix  
of weights  $\mathbf{w}_{ij}$

$$\sum_{i=1}^4 \sum_{j=1}^4 w_{ij} x_{i+p, j+q} + b$$

For neuron (p,q) in hidden  
layer


For a neuron in hidden layer:

- Take inputs from patch
- Compute weighted sum
- Apply bias

- 1) applying a window of weights
- 2) computing linear combinations
- 3) activating with non-linear function



# CNNs: Spatial Arrangement of Output Volume



**Layer Dimensions:**

$h \times w \times d$

**Stride:**



Filter Step Size

**Receptive Field:**

Locations in input image that a node is path-connected to.

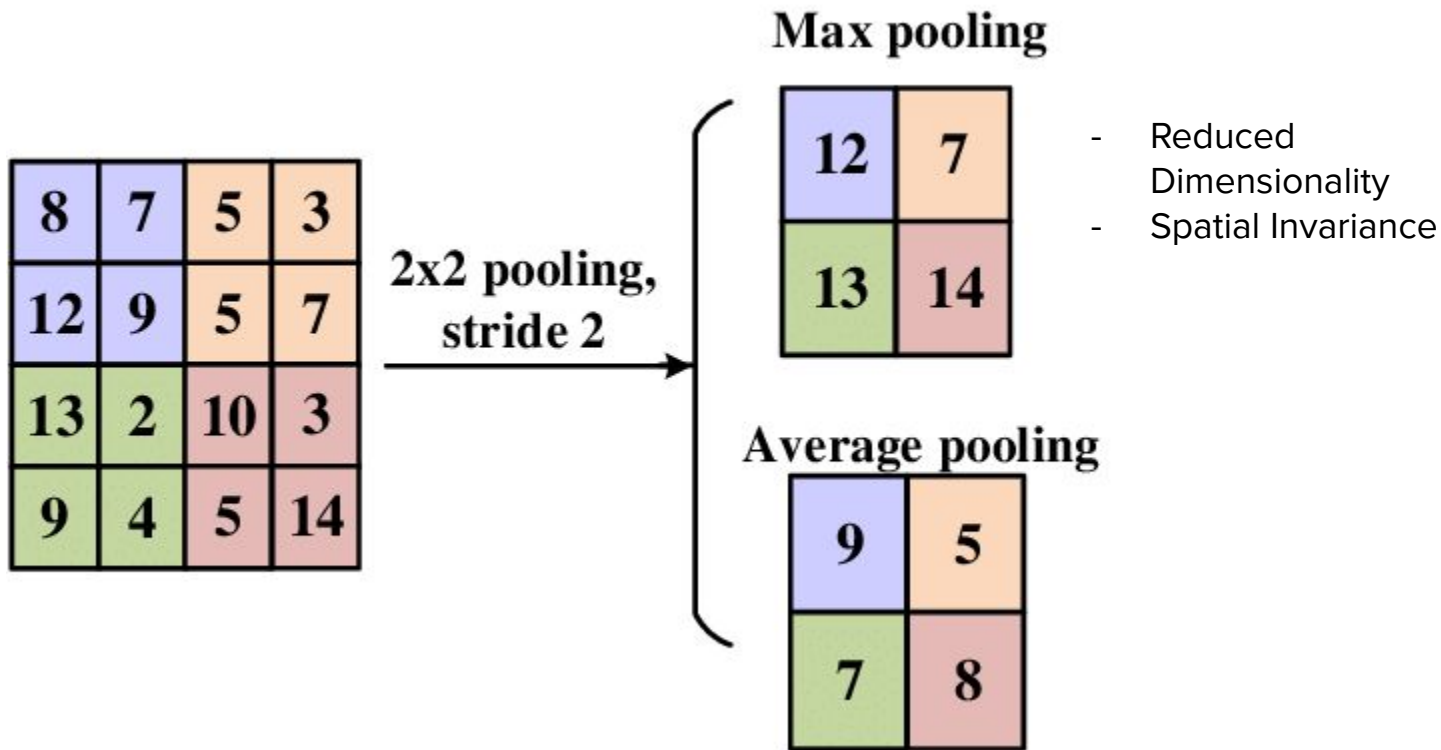
# Introducing Non-Linearity

- Apply after every convolution operation (i.e., after convolutional layers)
- ReLU: pixel-by-pixel operation that replaces all negative values by zero. **Non-linear operation!**




Rectified Linear Unit (ReLU)

# Pooling



# Representation Learning in Deep CNNs



Edges, dark spots

Eyes, ears, nose

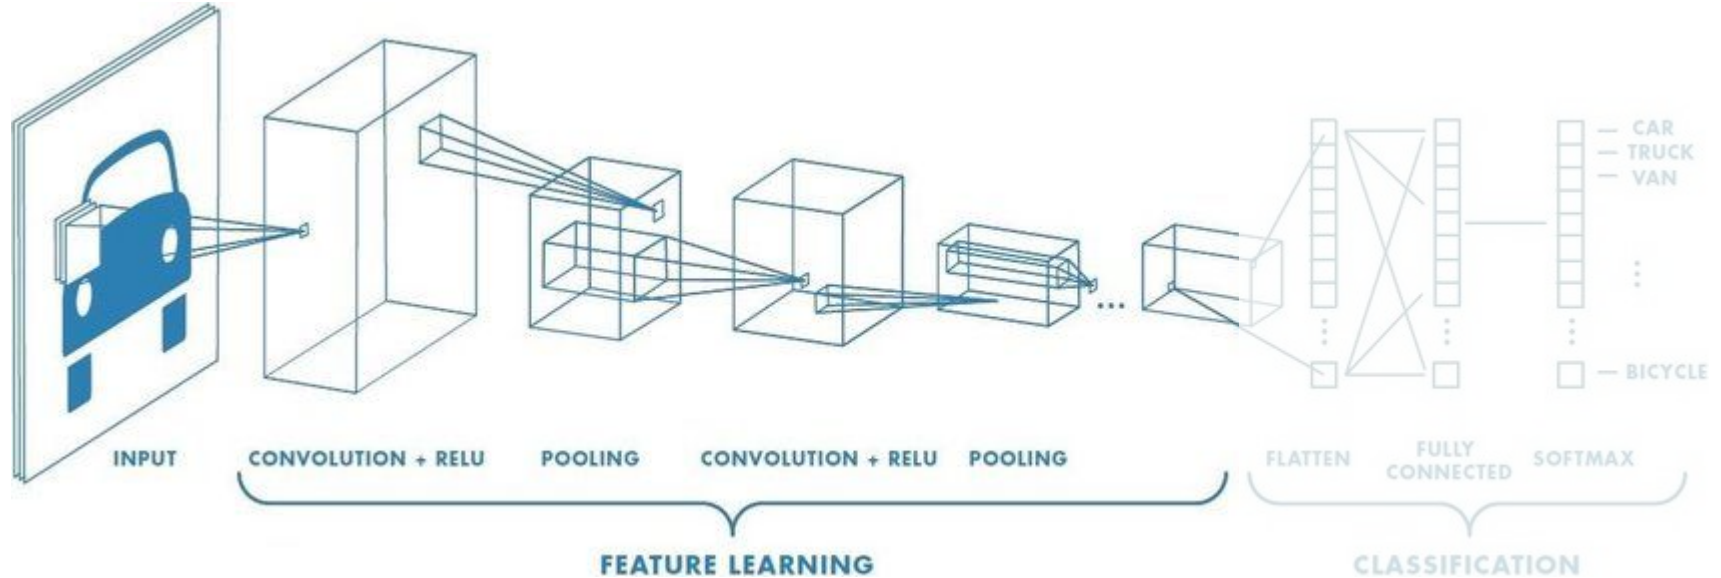
Facial Structure

Conv Layer 1

Conv Layer 2

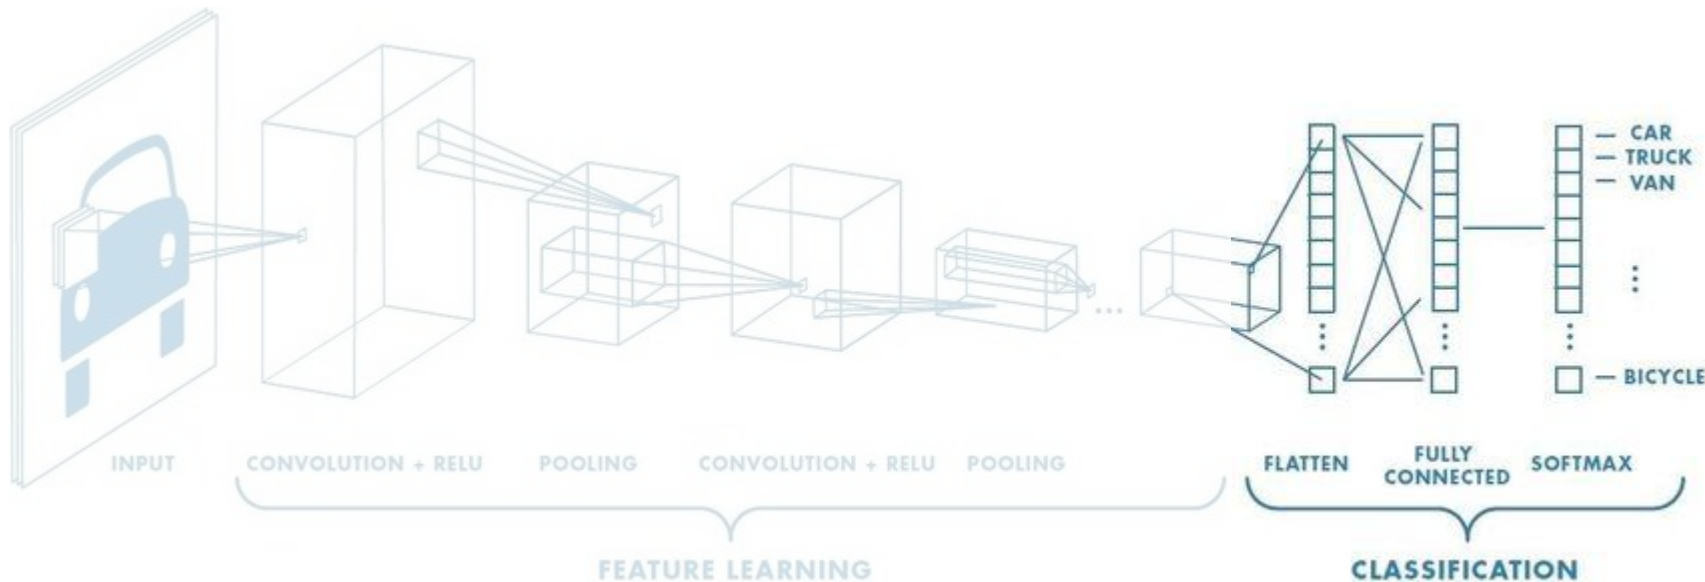
Conv Layer 3

# CNNs for Classification: Feature Learning



1. Learn features in input image through **convolution**
2. Introduce **non-linearity** through activation function (real-world data is non-linear!)
3. Reduce dimensionality and preserve spatial invariance with **pooling**

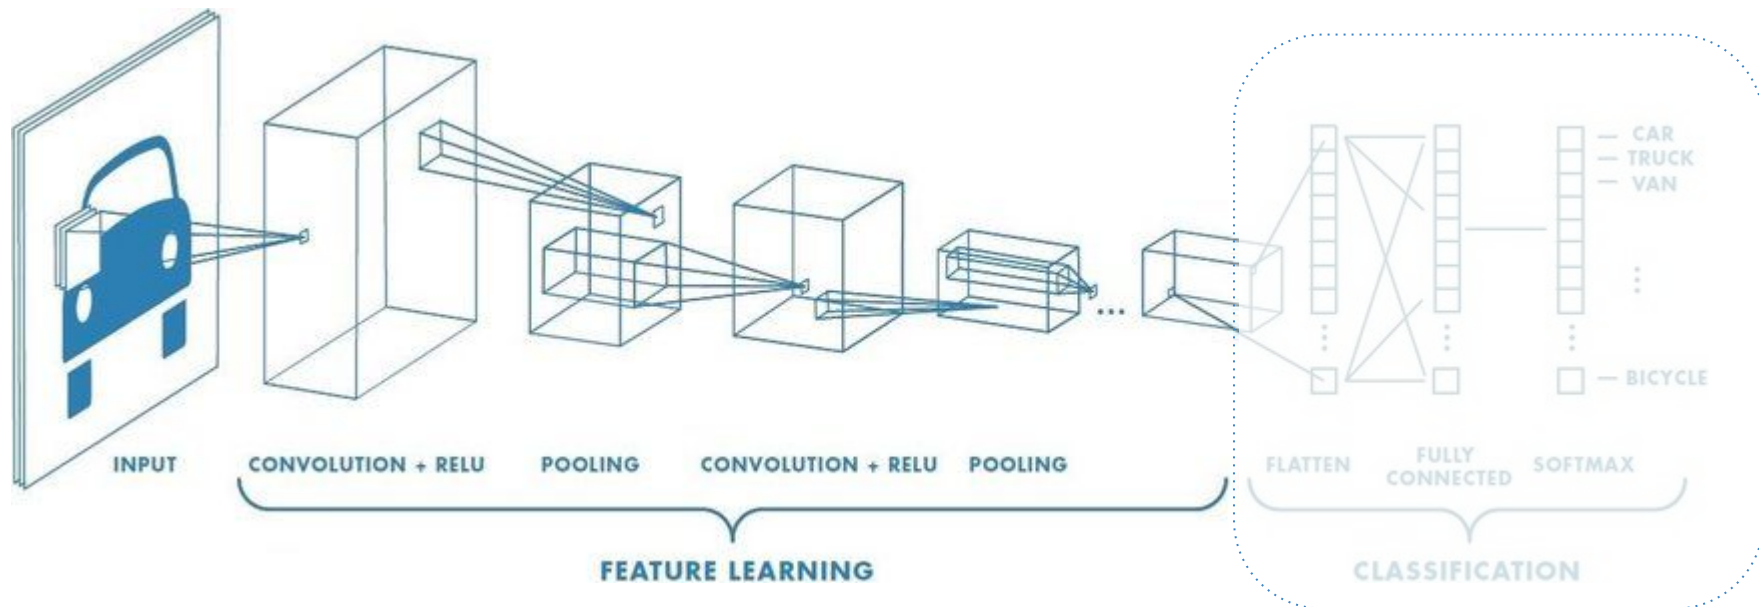
# CNNs for Classification: Class Probabilities



- CONV and POOL layers output high-level features of input
- Fully connected layer uses these features for classifying input image
- Express output as **probability** of image belonging to a particular class

$$\text{softmax}(y_i) = \frac{e^{y_i}}{\sum_j e^{y_j}}$$

# CNNs for Classification: Feature Learning



Classification  
Object Detection  
Segmentation  
Generative AI

# **Next Lecture: Different CNN Architectures, Transfer Learning!**